
The Effect of Instruction Padding on SFI Overhead

Navid Emamdoost

navid@cs.umn.edu

Stephen McCamant

mccamant@cs.umn.edu

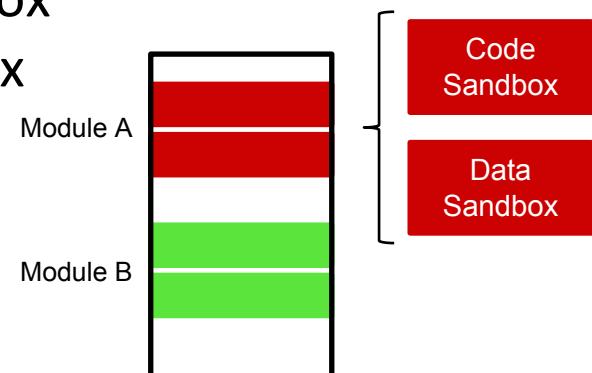


UNIVERSITY OF MINNESOTA

Driven to DiscoverSM

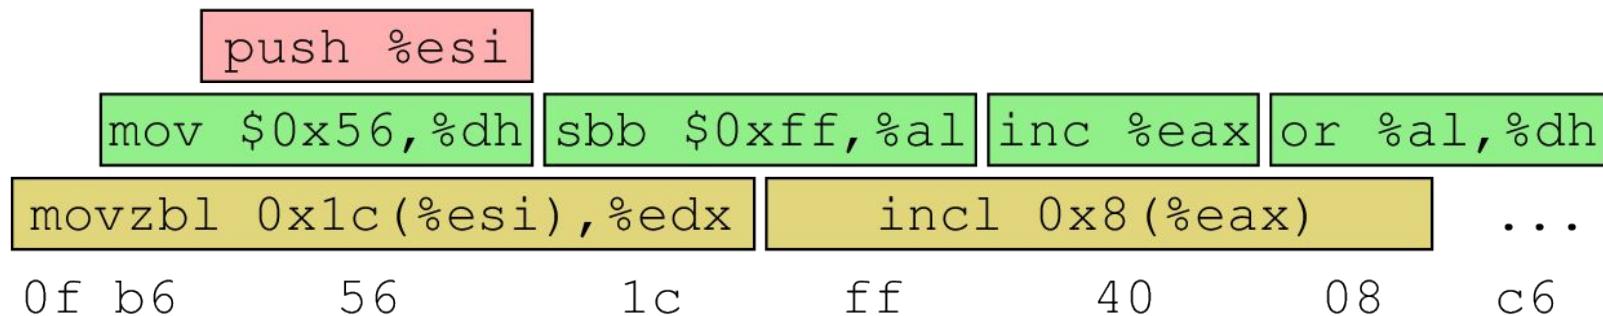
Software-based Fault Isolation

- Software-based isolation
 - Incorporate an untrusted module
 - All modules in same virtual address
 - Provide efficient communication
- Security Policy:
 - No code is executed outside the sandbox
 - No data is changed outside the sandbox



CISC Architectures

- Variable length instructions
- Processor can jump to any byte
- Hard to make hidden instructions safe
- Solution: Instruction Alignment



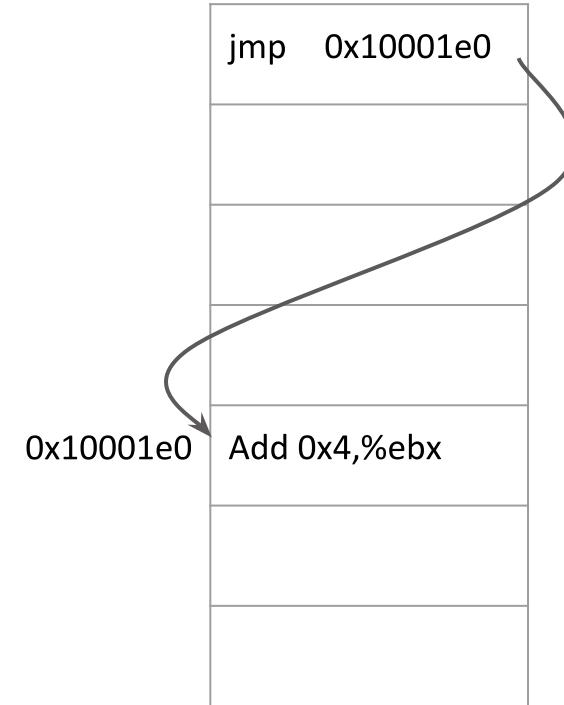


Our Focus

- Improving current isolation mechanism performance
 - Guaranteeing same level of security
- Native Client
 - Software-based Fault Isolation implementation for CISC architectures
 - Instruction padding to enforce address layout invariant
 - Padding imposes runtime overhead
 - Changed padding scheme
 - Updated validator

Unsafe Instructions

- Jump or store instructions
 - Change Control flow
 - Change data
- Addressing issue
 - `jmp 0x10001e0`
 - `mov %eax,0x11020028`



Unsafe Instructions

- Jump or store instructions
 - Change Control flow
 - Change data
- Addressing issue
 - `jmp 0x10001e0`
 - `jmp *%ecx`
 - `mov %eax,0x11020028`
 - `mov $0x1b80,(%ecx)`



Alignment in NaCl

- Divide memory into 32-byte *Bundles*
- Target of indirect jump is placed at the beginning of bundles
- No instruction is allowed to cross bundle boundary

```
...  
1060451: 83 c8 01          or    $0x1,%eax  
1060454: 88 41 28          mov    %al,0x28(%ecx)  
1060457: 8b 0e              mov    (%esi),%ecx  
1060459: 8b 51 04          mov    0x4(%ecx),%edx  
106045c: 85 d2              test   %edx,%edx  
106045e: 66 90              xchg   %ax,%ax  
  
1060460: 0f 88 9a 01 00 00  js    1060600  
1060466: 8b 01              mov    (%ecx),%eax  
1060468: 8d 3c 95 00 00 00 00 lea    0x0(%edx,4),%edi  
106046f: 89 d6              mov    %edx,%esi  
1060471: 83 ea 01          sub    $0x1,%edx  
1060474: 83 e6 07          and    $0x7,%esi  
1060477: 83 fa ff          cmp    $0xffffffff,%edx  
106047a: 8d b6 00 00 00 00  lea    0x0(%esi),%esi  
  
1060480: c7 04 38 84 06 03 11  movl   $0x11030684,(%eax,%edi,1)  
1060487: 8d 47 fc          lea    -0x4(%edi),%eax  
106048a: 0f 84 70 01 00 00  je    1060600  
...
```

Alignment in NaCl

- Divide memory into 32-byte *Bundles*
- Call instructions placed at the end of bundles
 - For the sake of return address alignment

```
10ee700: 55          push  %ebp
10ee701: 89 e5        mov    %esp,%ebp
10ee703: 83 ec 18     sub    $0x18,%esp
10ee706: c7 04 24 2c 5e 01 11  movl   $0x11015e2c,(%esp)
10ee70d: 8d b4 26 00 00 00 00  lea    0x0(%esi,%eiz,1),%esi
10ee714: 8d bc 27 00 00 00 00  lea    0x0(%edi,%eiz,1),%edi
10ee71b: e8 a0 72 f5 ff  call   10459c0

10ee720: 89 ec        mov    %ebp,%esp
10ee722: 5d          pop    %ebp
10ee723: 59          pop    %ecx
10ee724: 83 e1 e0     and    $0xfffffffffe0,%ecx
10ee727: ff e1        jmp    *%ecx
...

```



UNIVERSITY
OF MINNESOTA

Padding vs Performance

- Inserting artificial alignments
 - Necessary for protecting against hidden instructions
 - Extra instructions (NOP)
 - Hurt CPU cache and pre-fetch
- Impose performance overhead

CBI NaCl

- Change NaCl padding scheme
 - Less NOPs
 - To decrease performance overhead
- Multipass Validator
 - We must guarantee sandboxing policy enforcement
 - Appropriate changes in validator



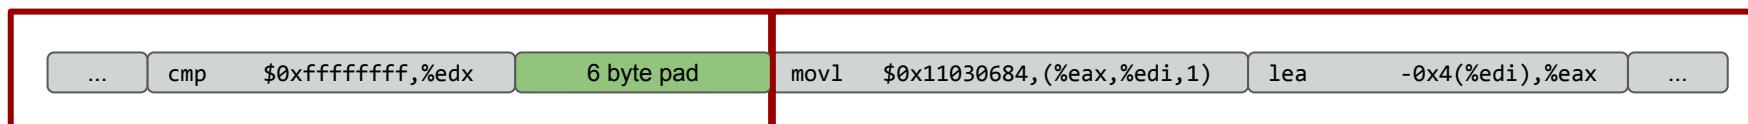
UNIVERSITY
OF MINNESOTA

Types of Padding

- Indirect jump target
 - Will be placed at the next bundle start
- Call instruction
 - Will be placed at the end of the bundle
- Cross bundle instruction
 - Will be pushed to the start of next bundle

Types of Padding

- Indirect jump target
 - Will be placed at the next bundle start
- Call instruction
 - Will be placed at the end of the bundle
- Cross bundle instruction
 - Will be pushed to the start of next bundle
 - Conservative



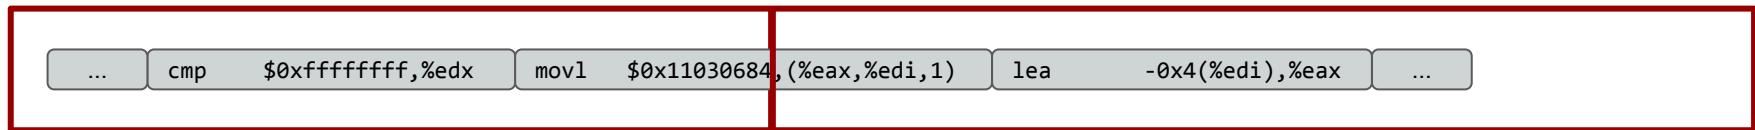
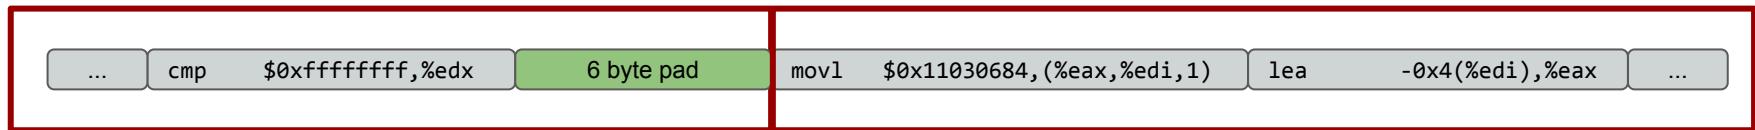


UNIVERSITY
OF MINNESOTA

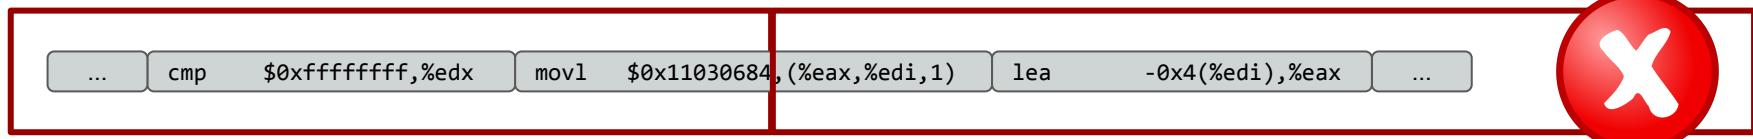
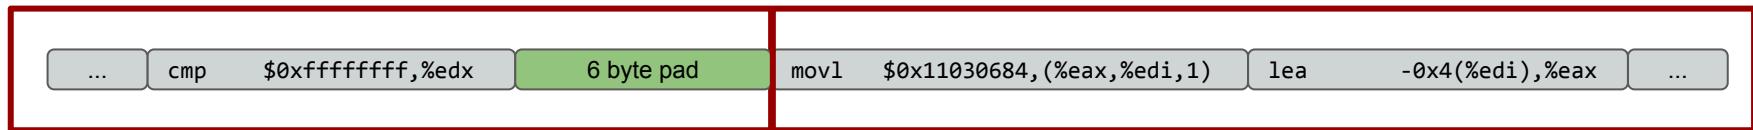
Pad Removal (simplified)

```
... cmp    $0xffffffff,%edx 6 byte pad movl  $0x11030684,(%eax,%edi,1) lea     -0x4(%edi),%eax ...
```

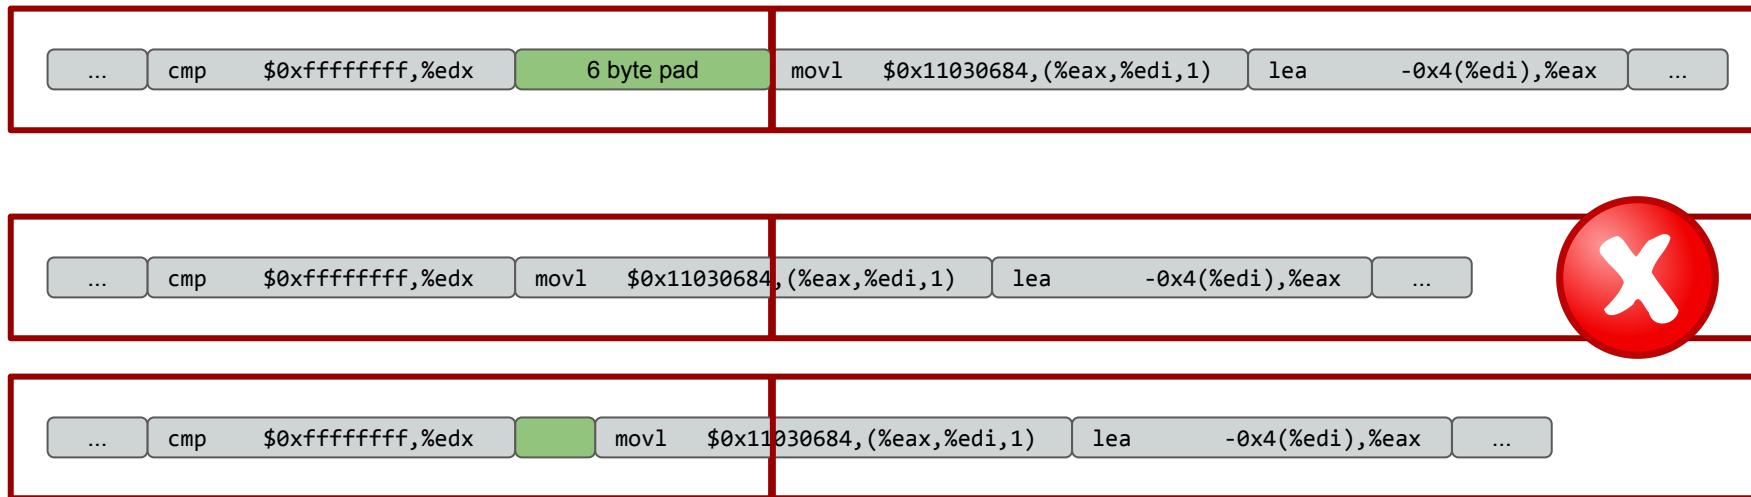
Pad Removal (simplified)



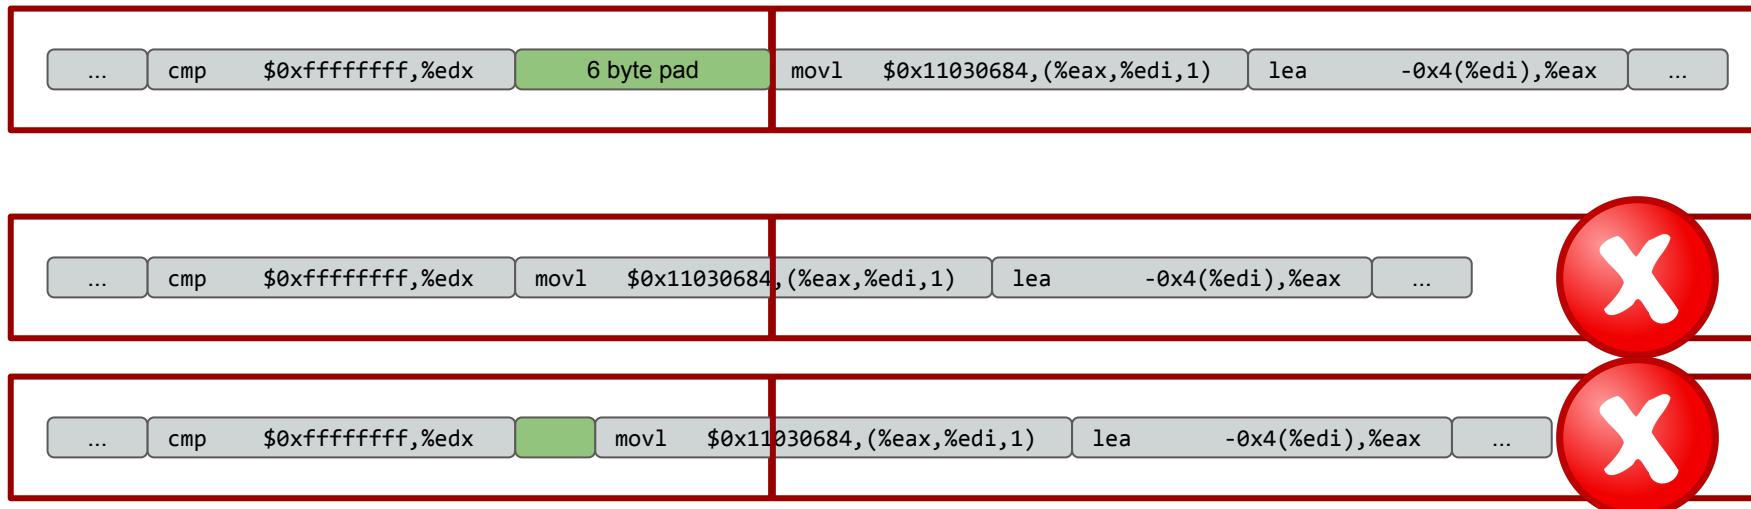
Pad Removal (simplified)



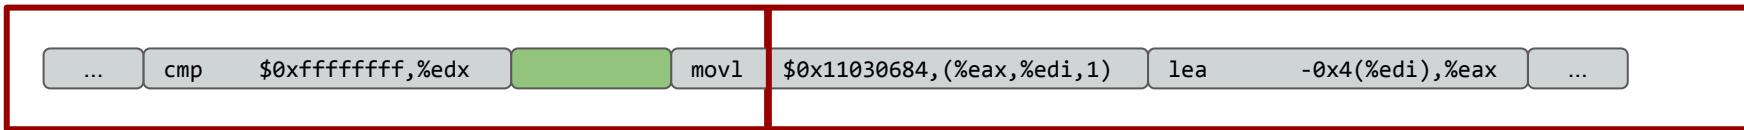
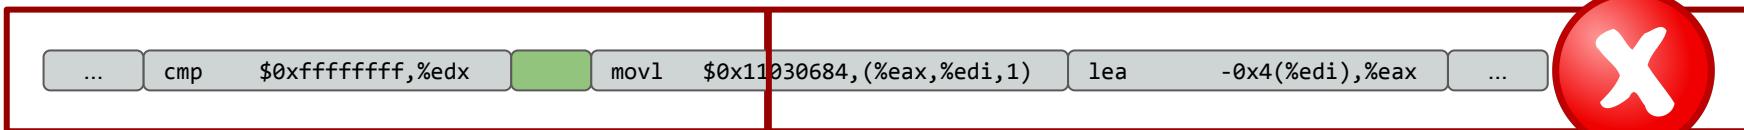
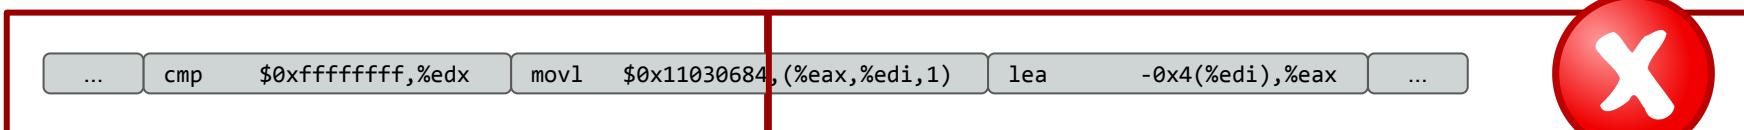
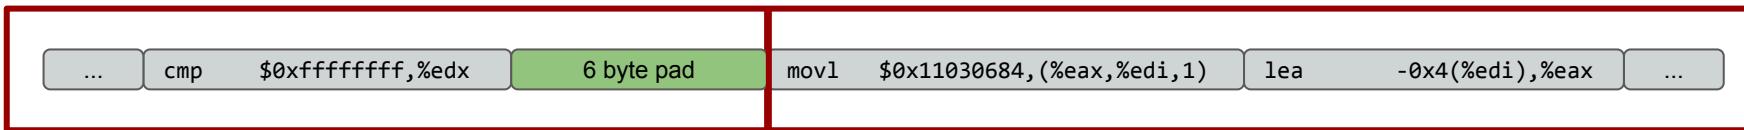
Pad Removal (simplified)



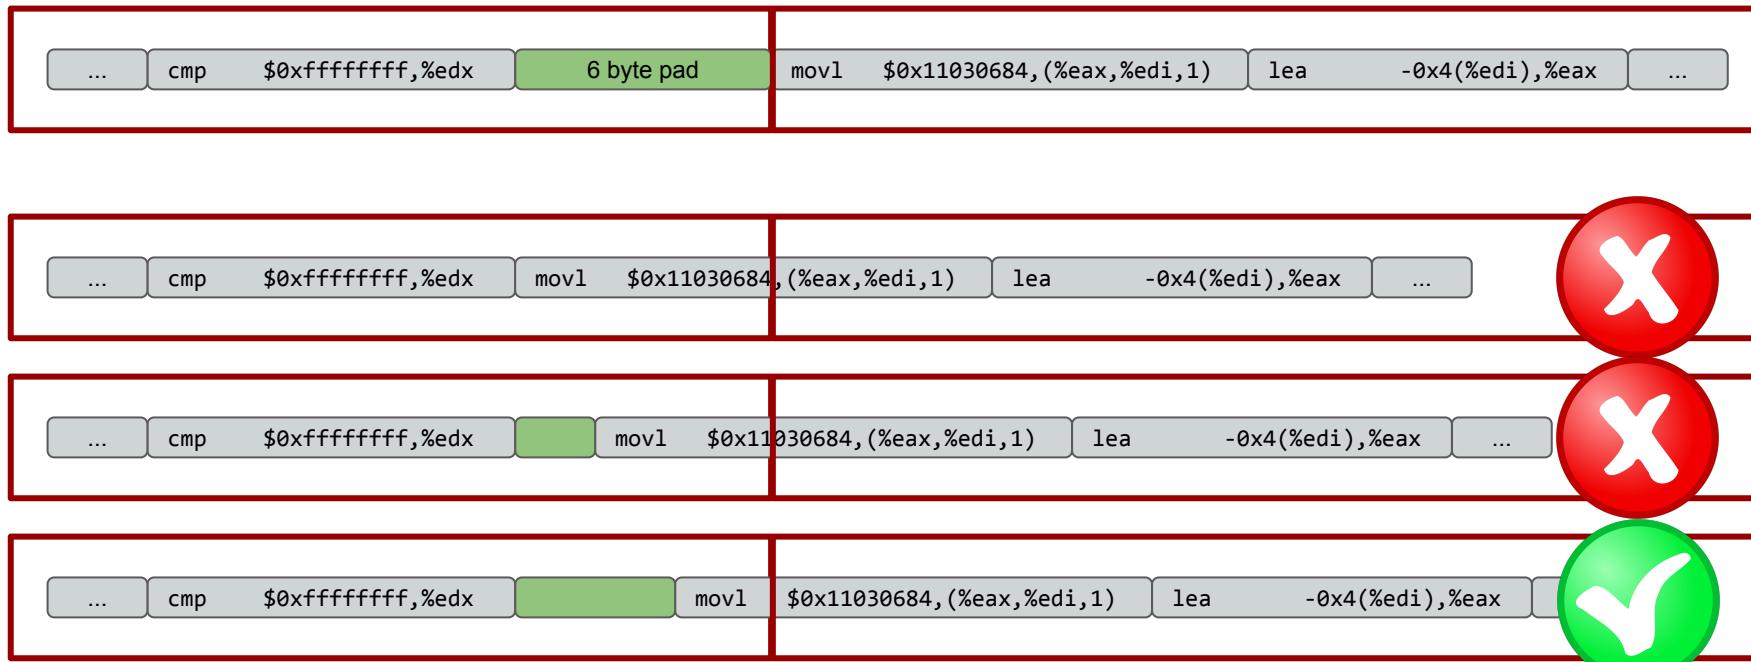
Pad Removal (simplified)



Pad Removal (simplified)



Pad Removal (simplified)





UNIVERSITY
OF MINNESOTA

NaCI Validator

- One pass: from the start to the end of code
- At each address checks the instruction
- For indirect branch
 - Checks presence of masking instruction
- For direct branch
 - Check the target is a valid instruction

Multipass Validator

- Challenge: Cross-Bundle Instructions

```
1060451: 83 c8 01      ...    or    $0x1,%eax
1060454: 88 41 28      mov    %al,0x28(%ecx)
1060457: 8b 0e          mov    (%esi),%ecx
1060459: 8b 51 04      mov    0x4(%ecx),%edx
106045c: 85 d2          test   %edx,%edx
106045e: 66 90          xchg   %ax,%ax
```

```
1060460: 0f 88 9a 01 00 00 js    1060600
1060466: 8b 01          mov    (%ecx),%eax
1060468: 8d 3c 95 00 00 00 lea   0x0(%edx,4),%edi
106046f: 89 d6          mov    %edx,%esi
1060471: 83 ea 01      sub    $0x1,%edx
1060474: 83 e6 07      and    $0x7,%esi
1060477: 83 fa ff      cmp    $0xffffffff,%edx
106047a: 8d b6 00 00 00 00 lea   0x0(%esi),%esi
```

```
1060480: c7 04 38 84 06 03 11  movl  $0x11030684,(%eax,%edi,1)
1060487: 8d 47 fc          lea    -0x4(%edi),%eax
106048a: 0f 84 70 01 00 00 je    1060600
...  
...
```

Multipass Validator

- Challenge: Cross-Bundle Instructions
- Multipass: start validation from every crossing point
 - Bundle start
 - Proceed to next if reached an already validated instruction



```
1060451: 83 c8 01 ... or $0x1,%eax
1060454: 88 41 28 mov %al,0x28(%ecx)
1060457: 8b 0e mov (%esi),%ecx
1060459: 8b 51 04 mov 0x4(%ecx),%edx
106045c: 85 d2 test %edx,%edx
106045e: 66 90 xchg %ax,%ax
```



```
1060460: 0f 88 9a 01 00 00 js 1060600
1060466: 8b 01 mov (%ecx),%eax
1060468: 8d 3c 95 00 00 00 00 lea 0x0(%edx,4),%edi
106046f: 89 d6 mov %edx,%esi
1060471: 83 ea 01 sub $0x1,%edx
1060474: 83 e6 07 and $0x7,%esi
1060477: 83 fa ff cmp $0xffffffff,%edx
106047a: 8d b6 00 00 00 00 lea 0x0(%esi),%esi
```



```
1060480: c7 04 38 84 06 03 11 movl
$0x11030684,(%eax,%edi,1)
1060487: 8d 47 fc lea -0x4(%edi),%eax
106048a: 0f 84 70 01 00 00 je 1060600
...
```



UNIVERSITY
OF MINNESOTA

Multipass Validator Correctness

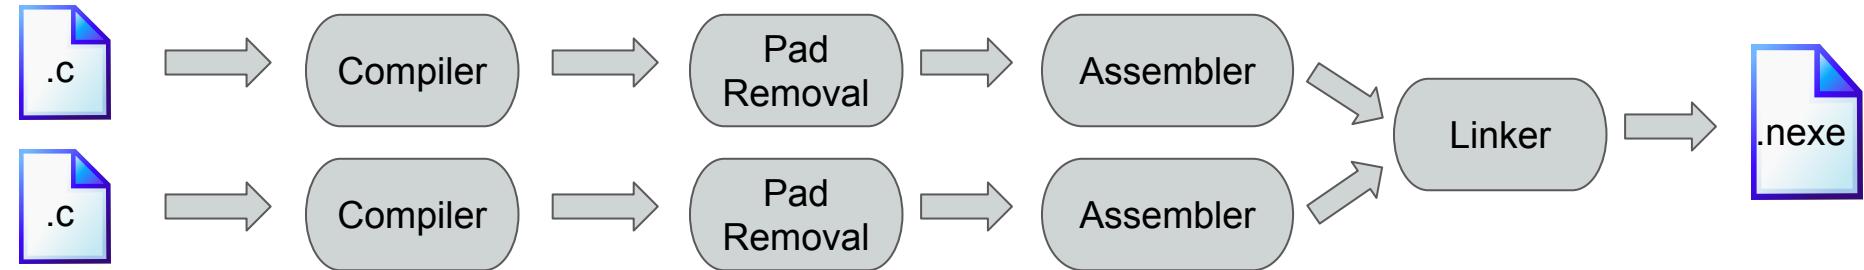
- NaCl sandboxing policy relies on:
 - Segment registers are not changed
 - Bytes in code segment are not altered after validation
- Multipass validator inherits NaCl validator instruction filtering
- Validation from every bundle start ensures all possible instructions are checked

Proof of Correctness

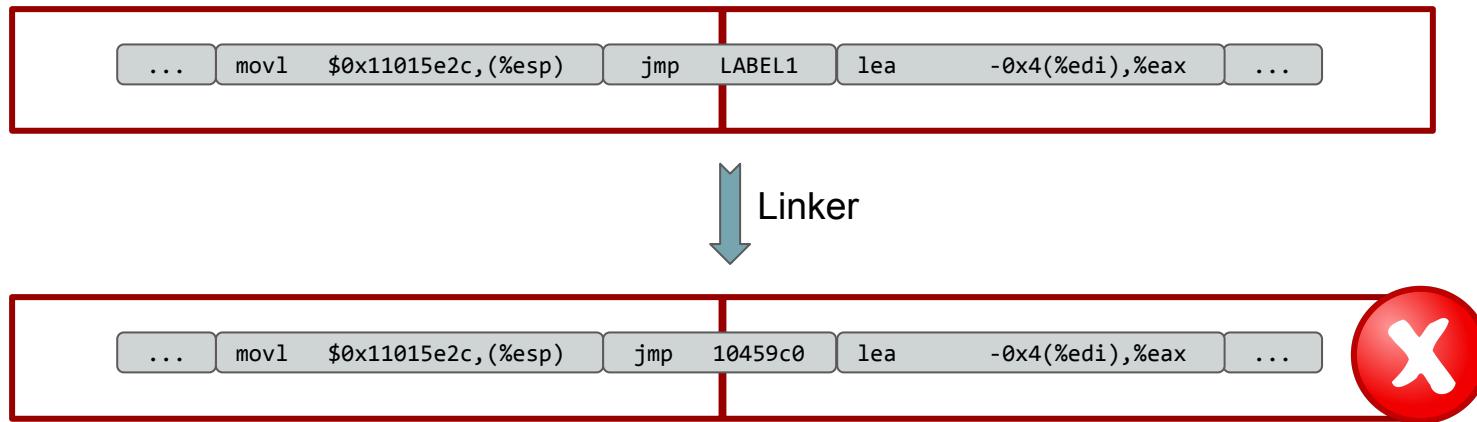
- Providing a formal proof for multipass validator correctness
 - Based on proof provided in RockSalt paper (PLDI'12)
 - Using Coq theorem prover
- RockSalt proof modeled a single-pass validator
- Modification to model multi-pass validator
 - Loops modeled as recursive functions
 - Inserted a recursive function to model outer loop
 - Starts validation from every bundle start
 - Reuse single-pass model as inner loop

Separate Compilation

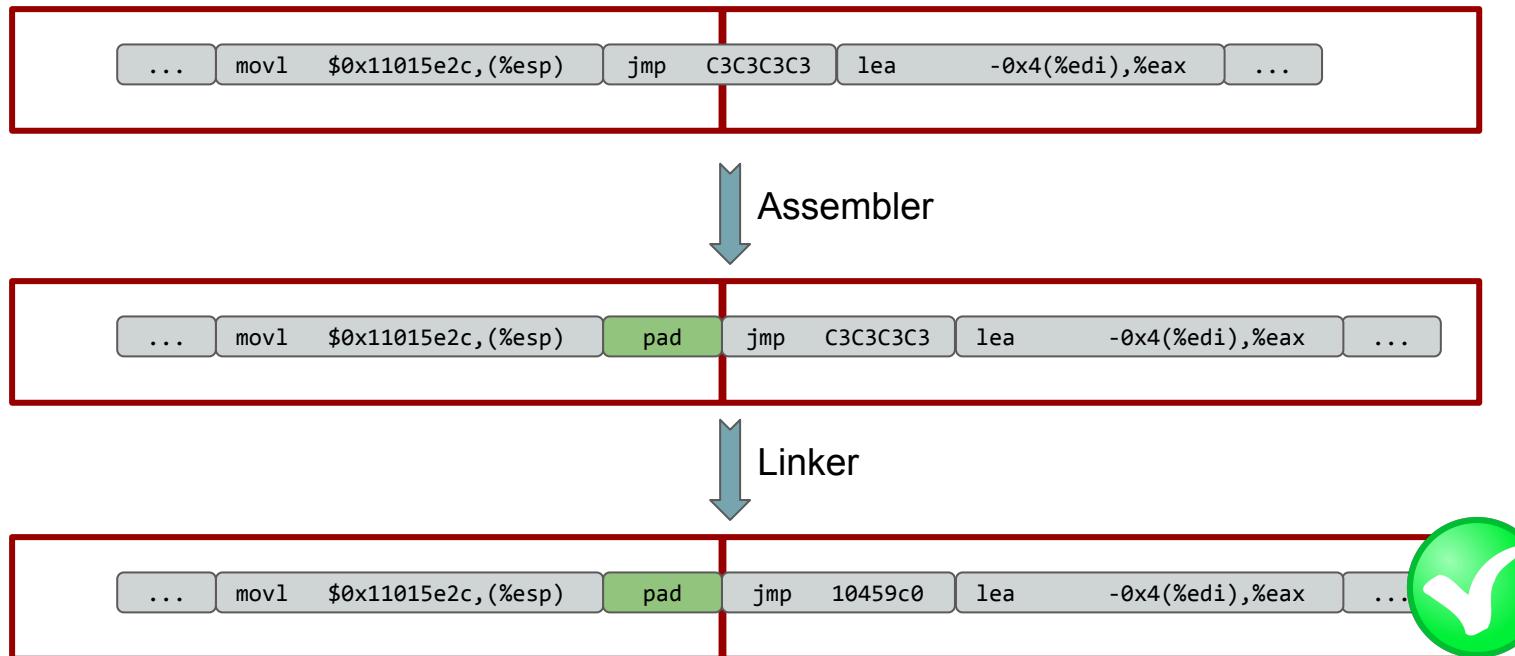
- Process each source file separately
 - Decide about the paddings to be removed
 - Assemble them into object files (using modified GAS)
 - Then link them together



Relocations Problem



Relocations Problem





More on compilation

- For some benchmarks $0xc3$ filtering was not enough
 - *Vortex* and *Crafty* from SPEC2000
 - *gcc*, *gobmk*, and *h264ref* from SPEC2006
- Reason: special values in relocation
 - Change the interpreted instruction length
- Caused the final binary fail validation
 - Even though with valid intermediate object files



UNIVERSITY
OF MINNESOTA

Example

- In nexe file:

```
66 89 81 9f 19 02 11 7f 99 e9 54 71 ff ff
```

- In object file:

```
66 89 81 c3 c3 c3 c3 7f 99 e9 54 71 ff ff
```

Example

- In nexe file:

```
66 89 81 9f 19 02 11 7f 99 e9 54 71 ff ff
```

```
104c53e: 66 89 81 9f 19 02 11    mov    %ax,0x1102199f(%ecx)
104c545: 7f 99                  jg     104c4e0 <Option+0x8ea0>
104c547: e9 54 71 ff ff        jmp    10436a0 <Option+0x60>
```

- In object file:

```
66 89 81 c3 c3 c3 c3 7f 99 e9 54 71 ff ff
```

```
1009a5e: 66 89 81 c3 c3 c3 c3    mov    %ax,-0x3c3c3c3d(%ecx)
1009a65: 7f 99                  jg     1009a00 <Option+0x8ea0>
1009a67: e9 54 71 ff ff        jmp    1000bc0 <Option+0x60>
```

Example

- In nexe file:

```
66 89 81 9f 19 02 11 7f 99 e9 54 71 ff ff
```

```
104c540: 81 9f 19 02 11 7f 99      sbbl    $0x7154e999,0x7f110219(%edi)
104c547: e9 54 71
104c54a: ff                      (bad)
104c54b: ff                      .byte 0xff
```

- In object file:

```
66 89 81 c3 c3 c3 c3 7f 99 e9 54 71 ff ff
```

```
1009a60: 81 c3 c3 c3 c3 7f      add     $0x7fc3c3c3,%ebx
1009a66: 99                      cltd
1009a67: e9 54 71 ff ff          jmp    0xfffff7160
```

Example

- In nexe file:

```
66 89 81 9f 19 02 11 7f 99 e9 54 71 ff ff
```



```
104c540: 81 9f 19 02 11 7f 99      sbbl    $0x7154e999,0x7f110219(%edi)
104c547: e9 54 71
104c54a: ff                      (bad)
104c54b: ff                      .byte 0xff
```

- In object file:

```
66 89 81 c3 c3 c3 c3 7f 99 e9 54 71 ff ff
```



```
1009a60: 81 c3 c3 c3 c3 7f      add     $0x7fc3c3c3,%ebx
1009a66: 99
1009a67: e9 54 71 ff ff        cltd
                                  jmp    0xfffff7160
```

- We filter *9f c3 c3 c3*, as well

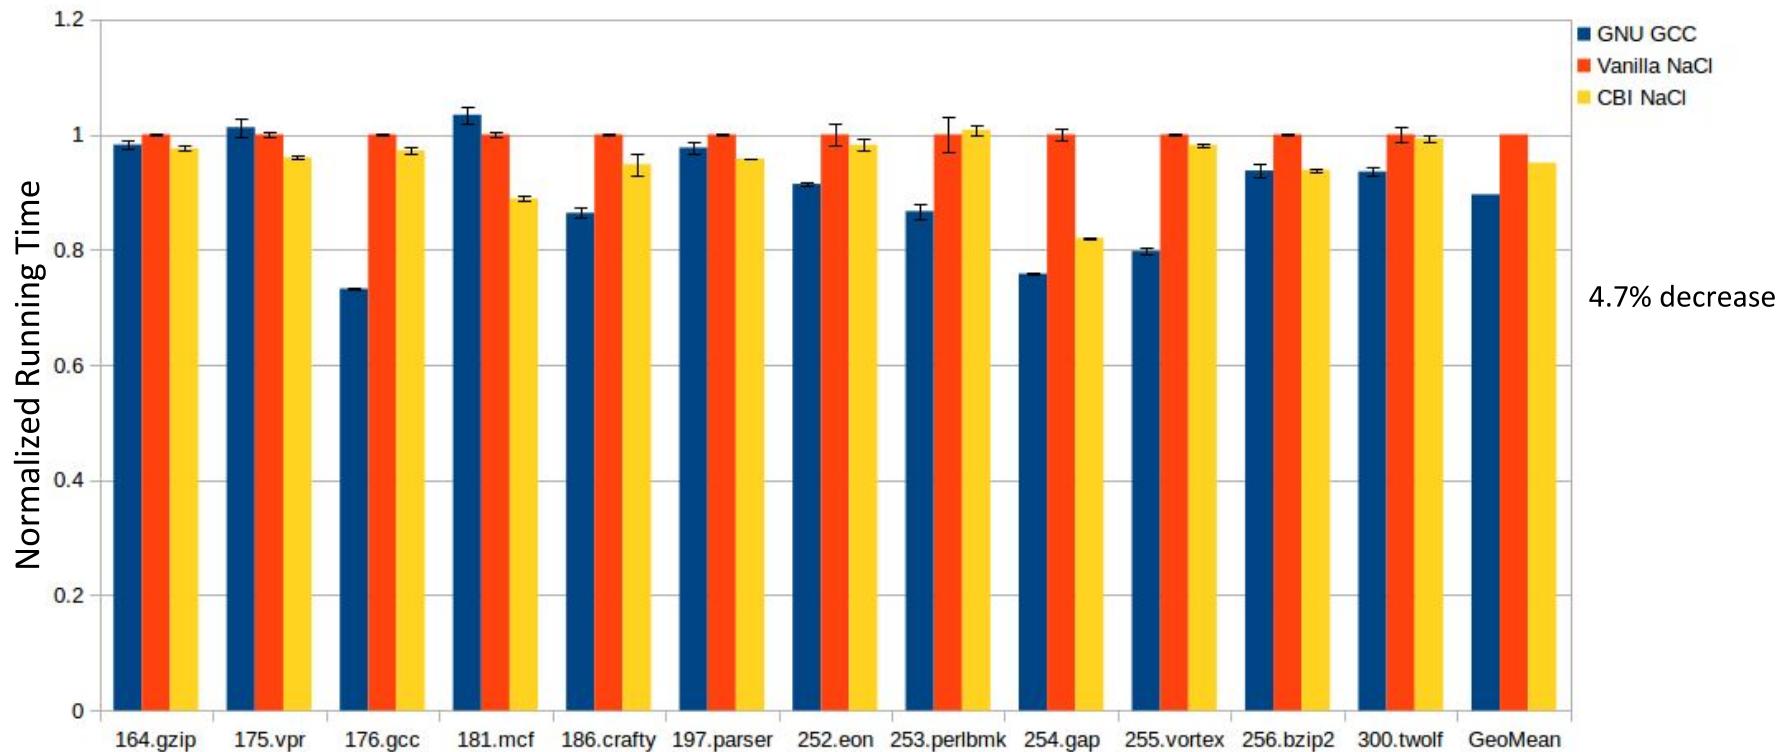


UNIVERSITY
OF MINNESOTA

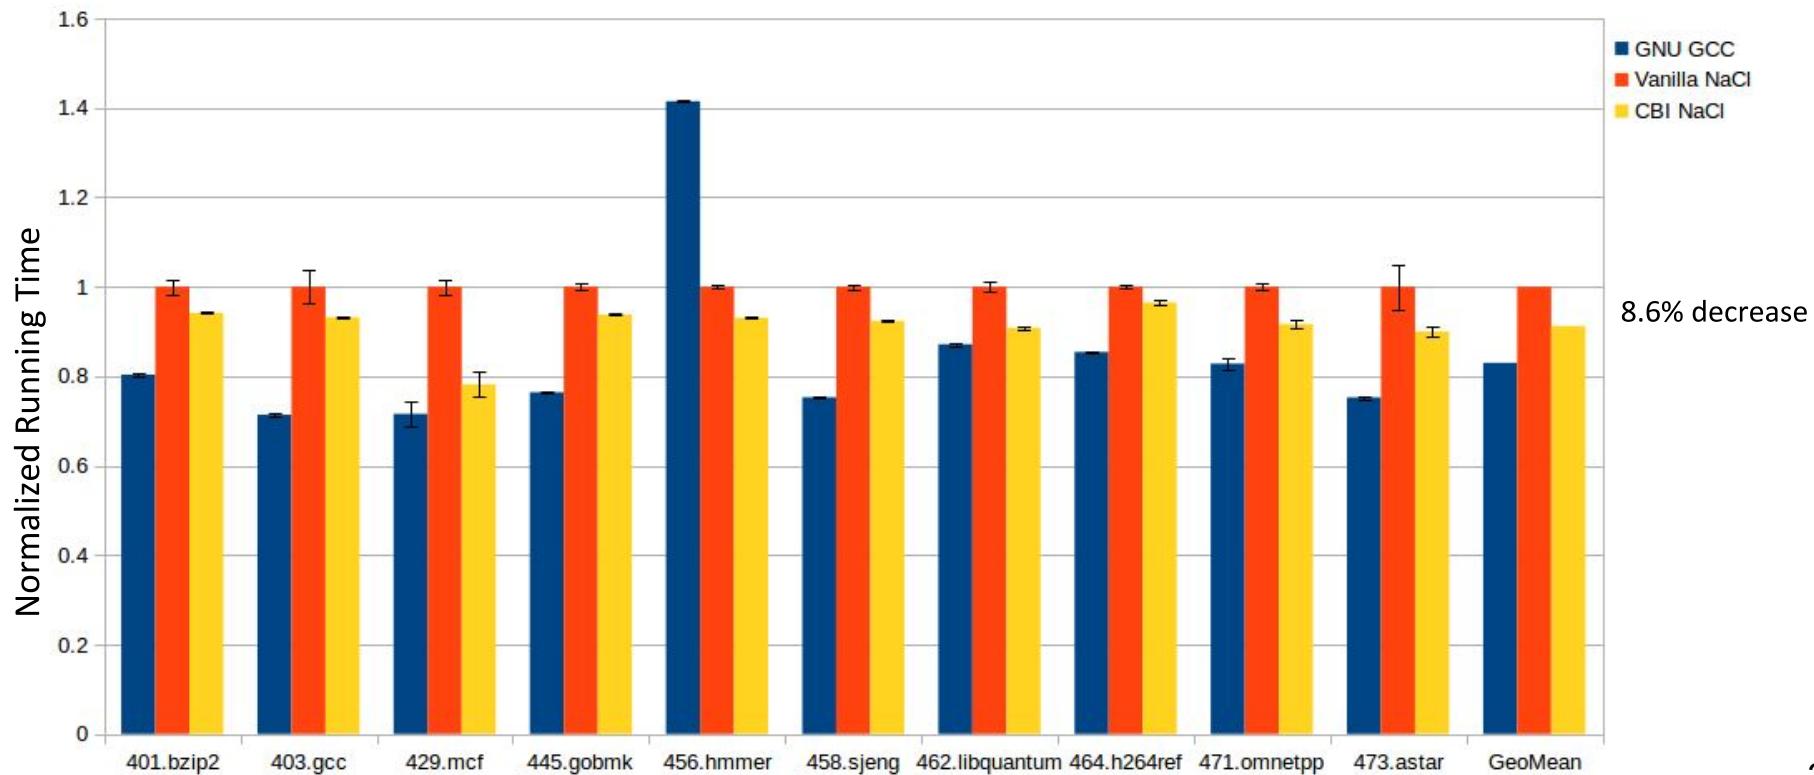
Performance Evaluation

- SPEC2000 and SPEC2006 benchmarks
- Built binaries using CBI NaCl build chain
 - SPEC2000: all 12 benchmarks built successfully
 - SPEC2006: 10 benchmarks were built successfully
 - One did not build even with GNU g++ 4.4.3
 - One fails to be executed in the NaCl
- Run for 11 times
- Discarding first executions

SPEC2000



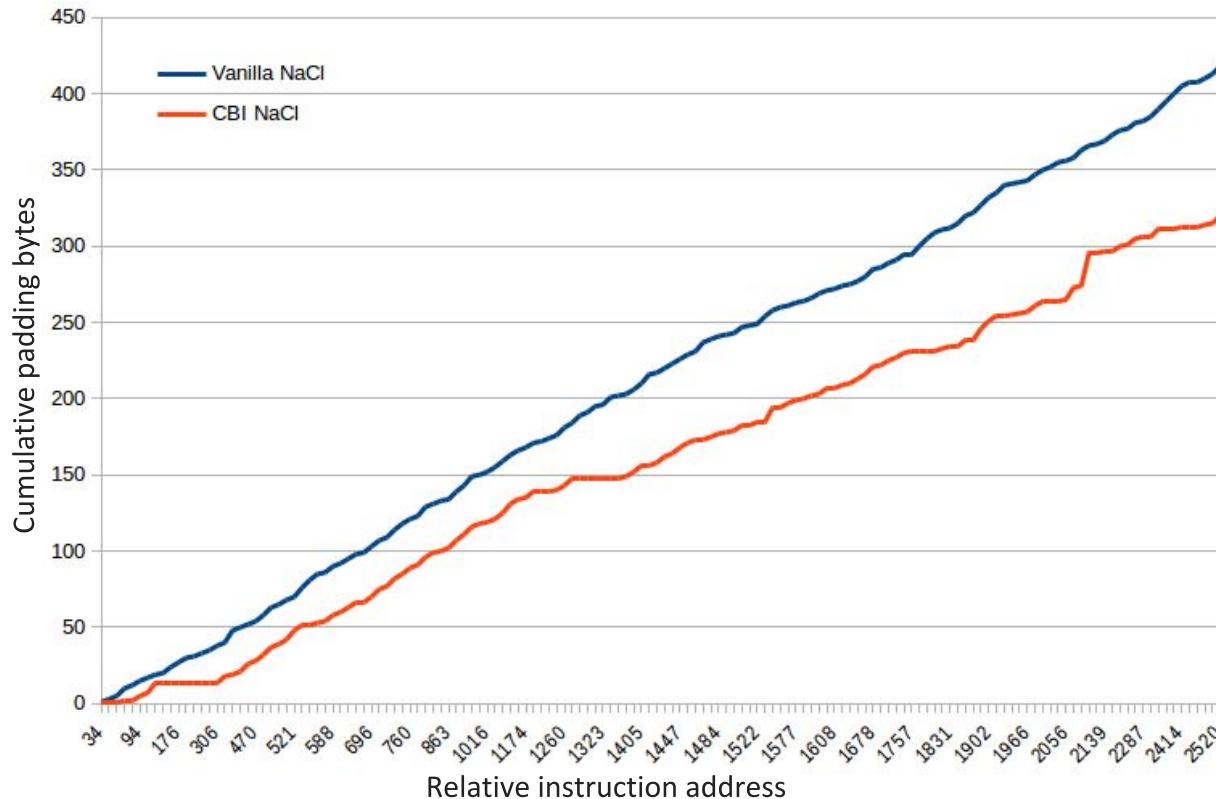
SPEC2006



Number of instructions

- We are removing NOPs
- Decrease in the number of instructions executed
 - 1.3% for SPEC2000 (collected via Valgrind)
 - 15% for SPEC2006 (collected via OProfile)
- We were able to remove more NOPs from SPEC2006
 - Less number of NOPs executed yields better performance

Pad Removal Success





On x64

- CBI NaCl is also applicable to x64
- x64 validation rules are more restrictive
 - No address segmentation
 - Restricted set of addressing are allowed
 - As a result: fewer overlapping instructions allowed
- CBI NaCl on x64
 - Rarely able to remove padding
 - Less than 0.5% performance improvement



UNIVERSITY
OF MINNESOTA

Future Work

- Replacing the greedy pad removal with a dynamic programming one
- Instruction re-ordering
 - To avoid branch prediction disruption
 - Reuse padded space to place actual instruction



Conclusion

- Effect of instruction level padding on SFI performance
- Change Native Client padding scheme
- Change NaCl build chain
 - Generate valid binaries
- Update validator appropriately and prove the correctness
- More than 8% performance improvement on average
- Project artifacts at:
 - <https://www-users.cs.umn.edu/~emamd001/cbi-nacl.html>



UNIVERSITY
OF MINNESOTA

Thank you for your attention
Q/A